# Support Vector Machines

Szu-Chi Chung

Department of Applied Mathematics, National Sun Yat-sen University

Generative and Discriminative Models for Classification

Discriminative   Generative

▸ Model the separating criterion $Pr(Y|X)$ directly

▸ Model the distribution of X in each of the classes $Pr(X|Y)$ separately, and then use Bayes theorem to flip things around and obtain $Pr(Y|X)$

Support Vector Machines

▶ The support vector machine is a generalization of a simple and intuitive classifier called the *maximal margin classifier* and *support vector classifier*

▶ Here we approach the two-class classification problem in a direct way:

  ▶ We try and find a plane that separates the classes in feature space

▶ If we cannot, we get creative in two ways:

  ▶ We <u>soften</u> what we mean by "separates", and

 ▶ We <u>enrich and enlarge the feature space</u> so that separation is possible
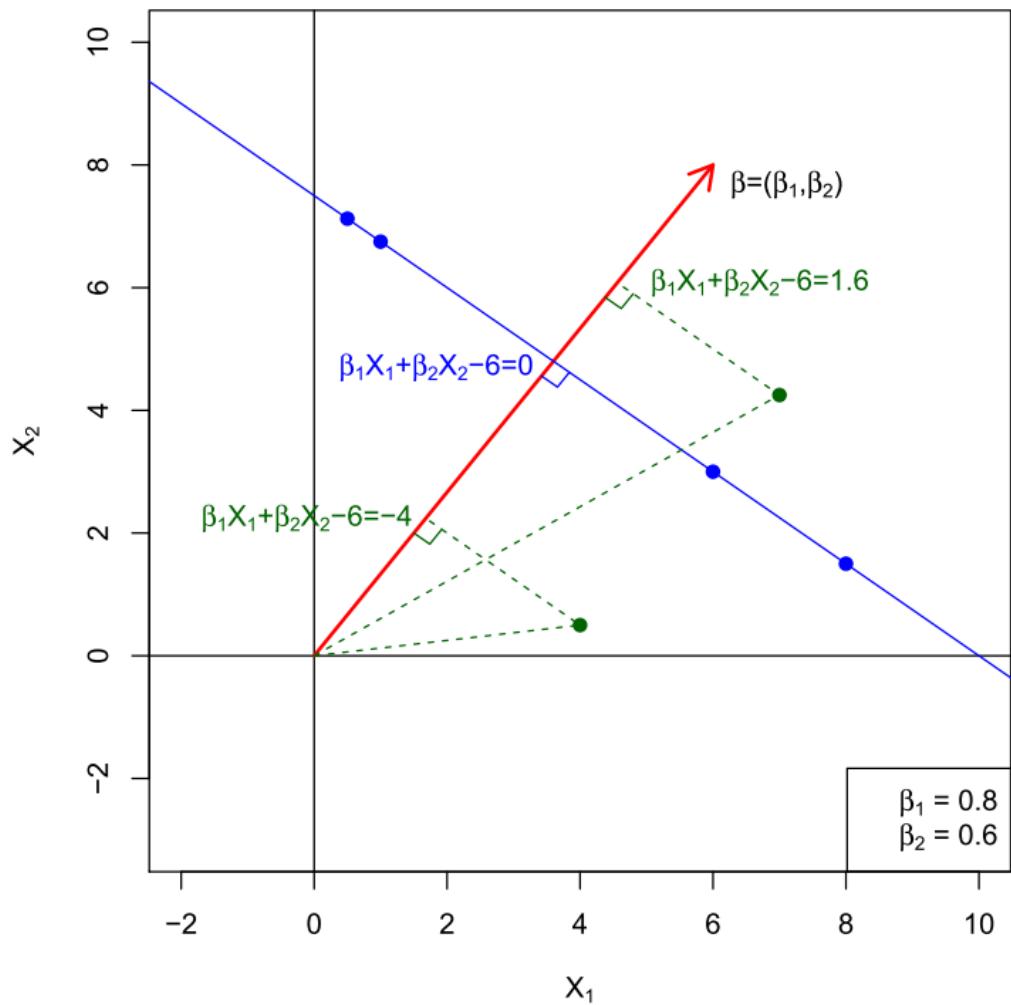
Why Support Vector Machine?

▶ For classification task we only need to know the hyperplane. Why bother a model?

▶ Support Vector Machine has the following advantages

   ✓ Still effective in cases where number of dimensions is greater than the number of samples

   ✓ Uses a subset of training points in the decision function (called support vectors), so it is also *efficient in both execution time and memory*

   ✓ Versatile: different Kernel functions can be specified for the decision function. Common kernels are provided, but it is also possible to specify custom kernels

x The disadvantage is that it does not directly provide probability estimates. In addition, if $p \gg n$, the regularization is crucial

What is a Hyperplane?

▸ A <u>hyperplane</u> in p dimensions is a flat affine subspace of dimension $p - 1$

  ▸ In general the equation for a hyperplane has the form
  $$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p = 0$$

  ▸ If $p = 2$ dimensions a hyperplane is a line

  ▸ If $p = 3$ dimensions a hyperplane is a plane

▸ If $\beta_0 = 0$, the hyperplane goes through the origin, otherwise not

▸ The vector $\beta = (\beta_1, \beta_2, \dots, \beta_p)$ is called the <u>normal vector</u> — it points in a direction orthogonal to the surface of a hyperplane
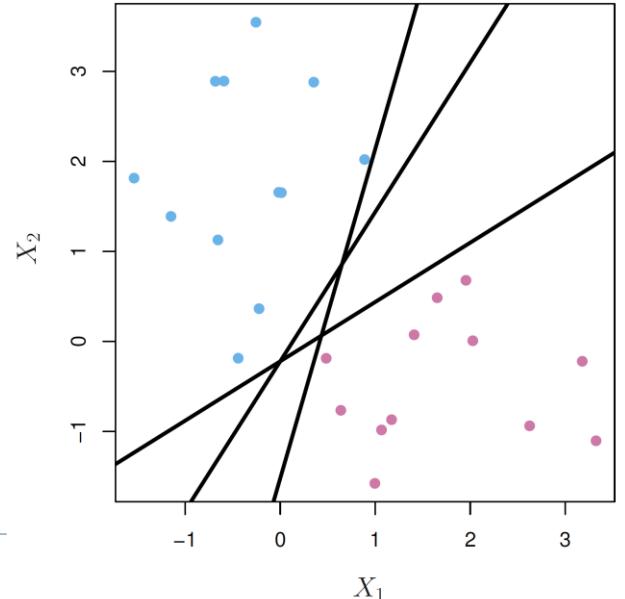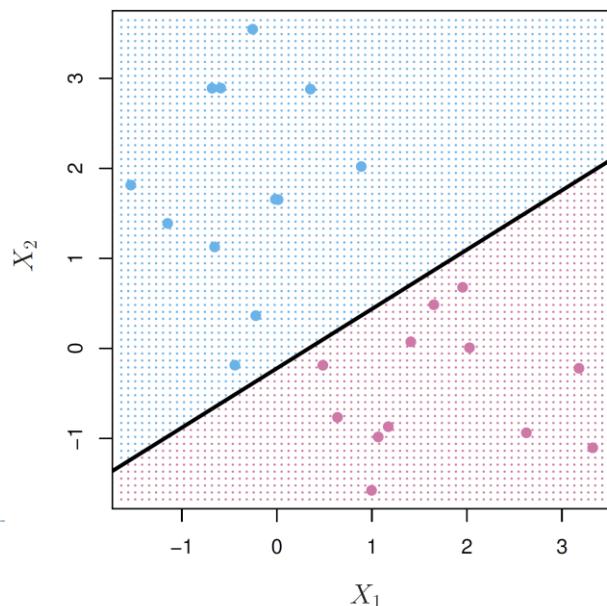
Hyperplane in 2 Dimensions

Note the distance between point (x_0, y_0) to line $\beta_1 X_1 + \beta_2 X_2 + \beta_0 = 0$ is $\dfrac{|\beta_1 x_0 + \beta_2 y_0 + \beta_0|}{\sqrt{\beta_1^2 + \beta_2^2}}$

, where $\dfrac{\beta_1 x_0 + \beta_2 y_0 + \beta_0}{\sqrt{\beta_1^2 + \beta_2^2}}$ is the scalar projection
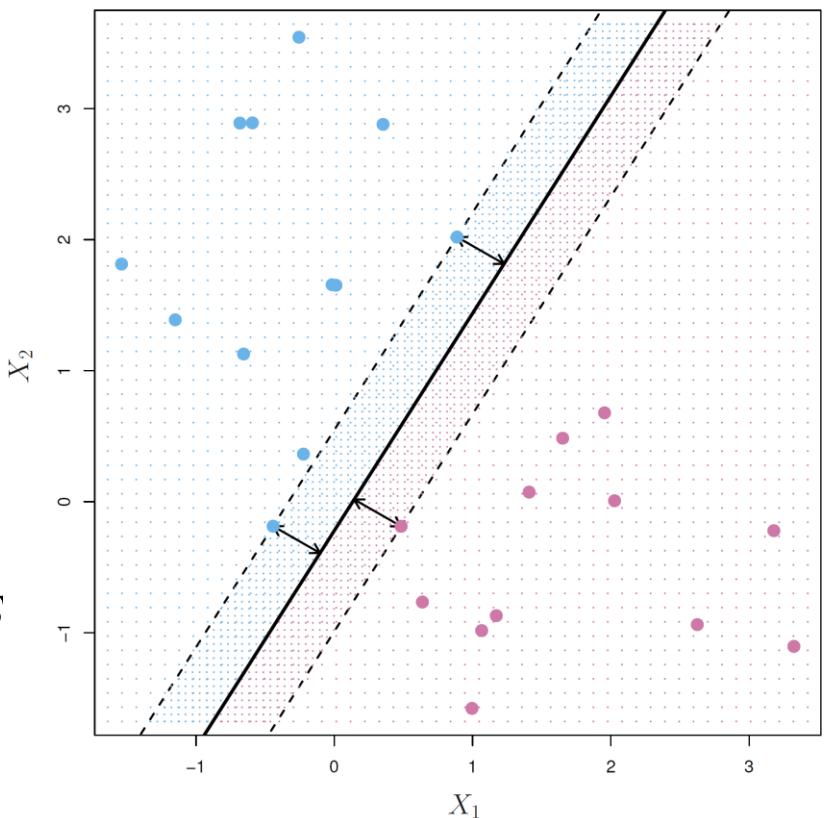
Separating Hyperplanes

▸ If $f(X) = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p$

  ▸ $f(X) > 0$ for points on one side of the hyperplane, and $f(X) < 0$ for points on the other

  ▸ If we code the colored points as $y_i = +1$ for blue, say, and $y_i = -1$ for purple, $f(X) = 0$ defines a separating hyperplane. Therefore, $y_i(\beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip}) > 0$ for all y_i, $i = 1, \ldots, n$!

  ▸ $y_i(\beta_0 + \beta_1 x_1^* + \cdots + \beta_p x_p^*)$ represents the confidence of our assignment

1. Maximal Margin Classifier

▸ Among all separating hyperplanes, find the one that makes the biggest *margin* between the two classes

  ▸ The minimal distance from the training observations to the hyperplane is known as the *margin*

  ▸ We try to choose the *maximal margin hyperplane*

  ▸ Classify a test observation based on which side of the maximal margin hyperplane it lies

▸ Three training observations are equidistant from the maximal margin which are known as *support vectors*

  ▸ The maximal margin hyperplane does not depends on all other training points

1. Maximal Margin Classifier
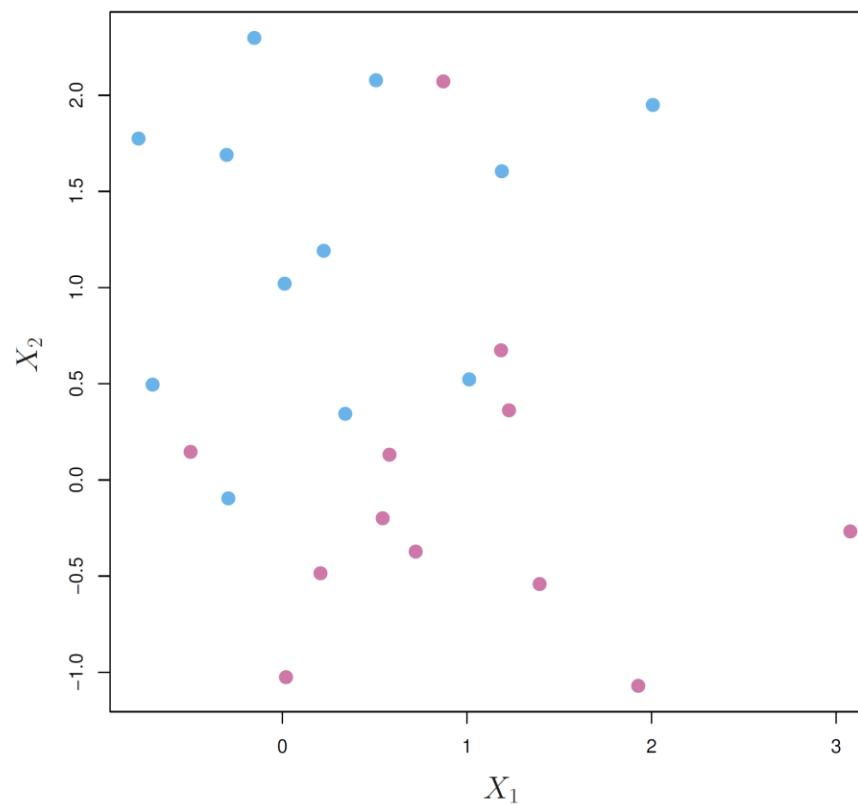
▸ The maximal margin hyperplane can be obtained via the constrained optimization problem

$$\max_{\beta_0,\beta_1,\dots,\beta_p} M$$

$$\text{subject to } \sum_{j=1}^{p} \beta_j^2 = 1,$$

$$y_i(\beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip}) \geq M$$

$$\text{for all } i = 1, \dots, n$$

▸ The first constraint means that the perpendicular distance from the ith observation to the hyperplane is given by $y_i(\beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip})$

▸ The second one guarantees that each observation will be on the correct side of the hyperplane with some cushion, provided that M is positive
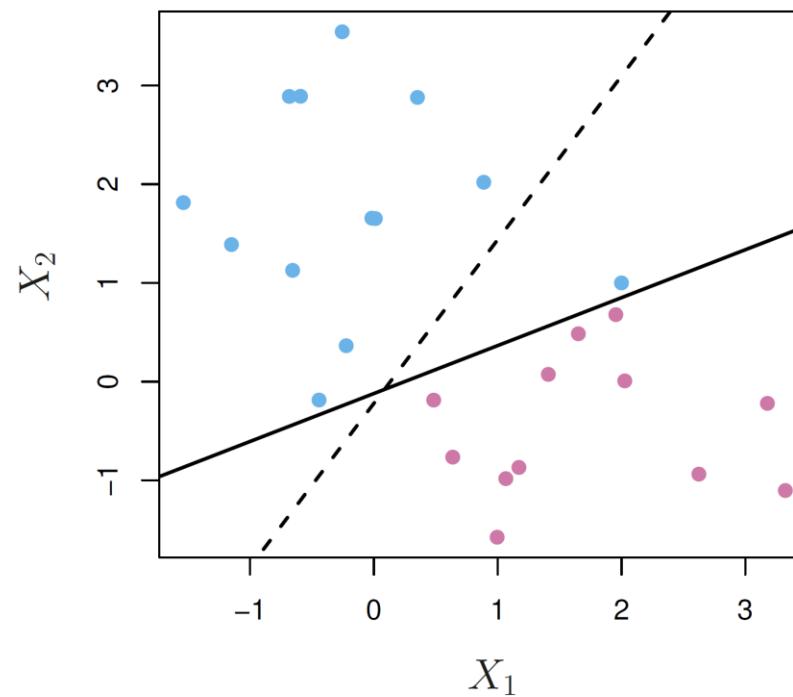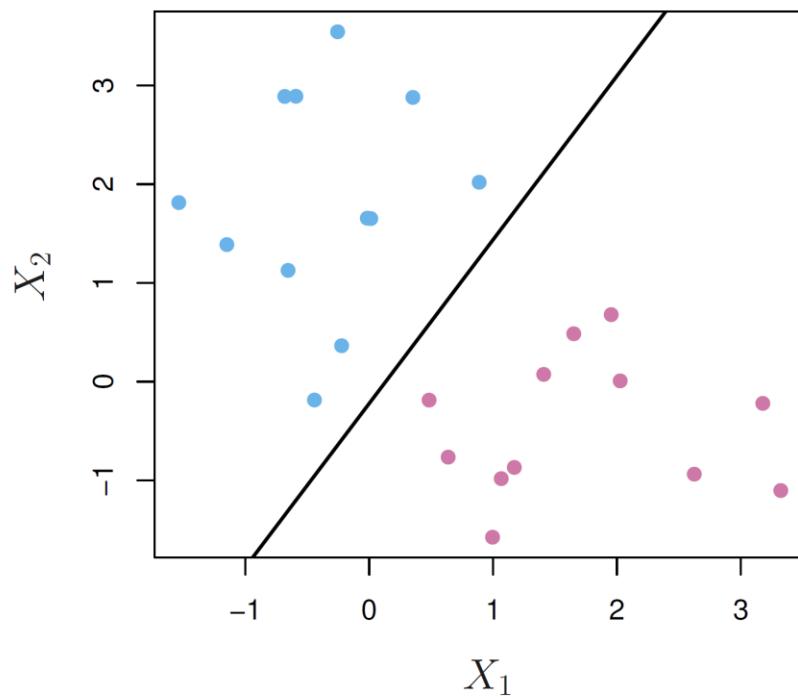
Non-separable Data

▸ The data below are not separable by a linear boundary

▸ This is often the case, unless $n < p$
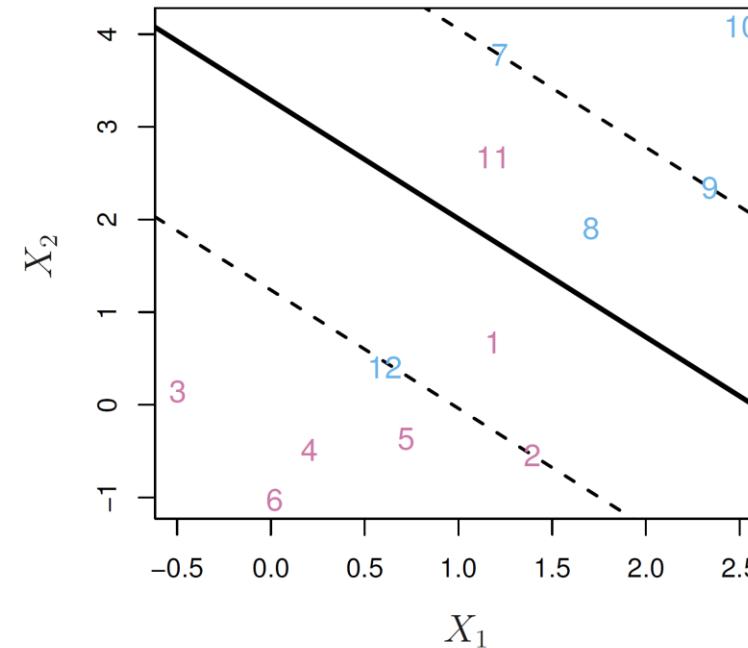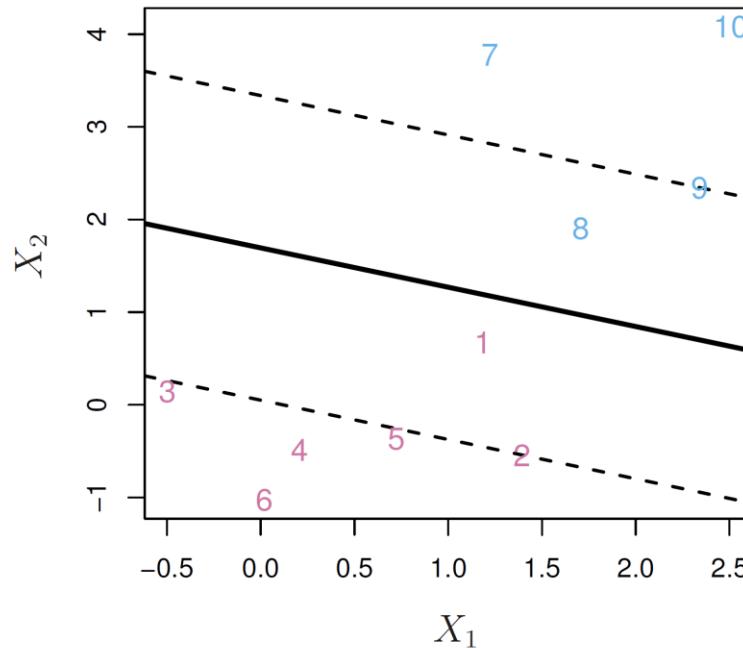
Noisy Data

▸ Sometimes the data are separable, but noisy. This can lead to a poor solution for the maximal-margin classifier

▸ The *support vector classifier* maximizes a *soft margin*

2. Support Vector Classifier (SVC)

▶ In this case, we might be willing to consider a *soft margin classifier* that does not perfectly separate the two classes, but

  ▸ Greater robustness to individual observations, and

  ▸ Better classification of *most* of the training observations

2. Support Vector Classifier (SVC)
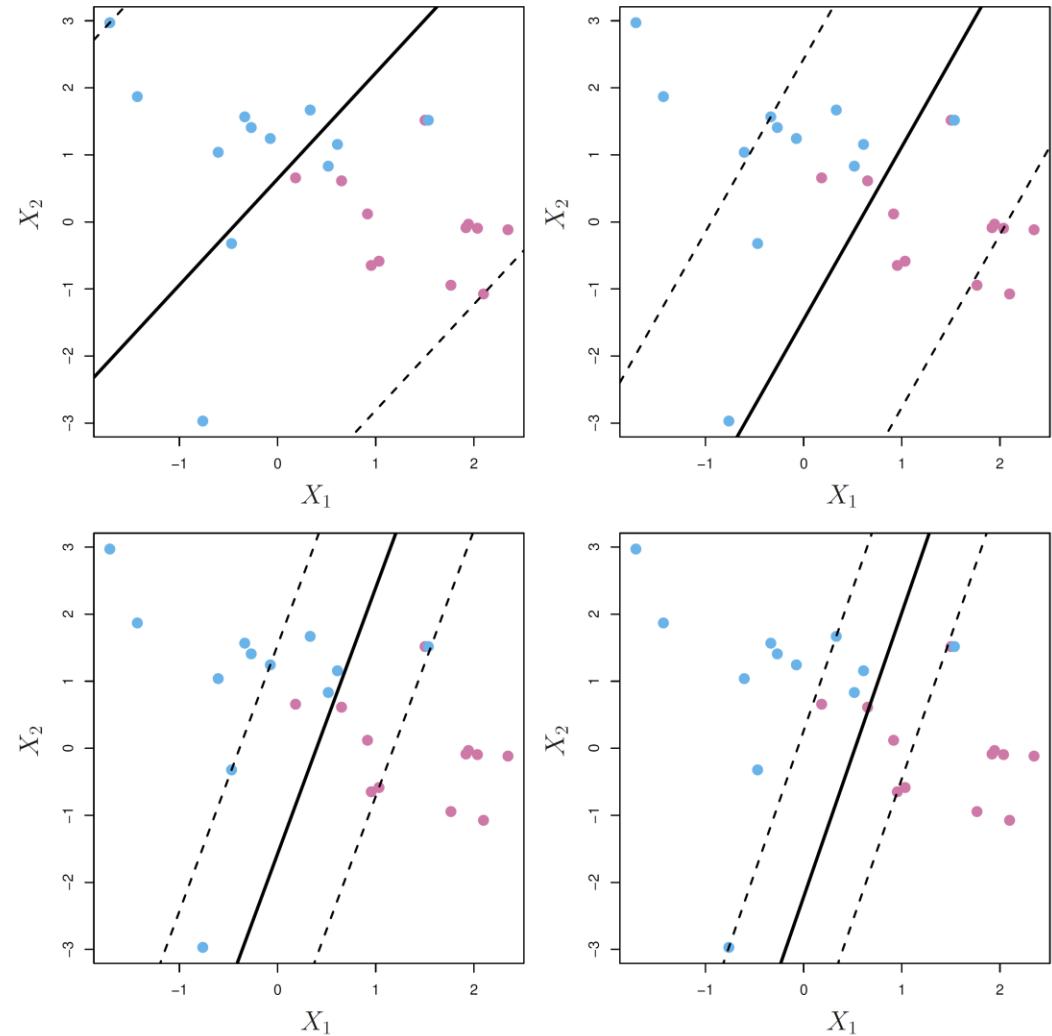
▸ The hyperplane is chosen to correctly separate most of the training observations into the two classes, but may misclassify a few observations via

$$\max_{\beta_0,\beta_1,\ldots,\beta_p,\epsilon_0,\epsilon_1,\ldots,\epsilon_n} M$$

$$\text{subject to } \sum_{j=1}^{p} \beta_j^2 = 1,$$

$$y_i\big(\beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip}\big) \geq M(1 - \epsilon_i),$$

$$\epsilon_i \geq 0, \sum_{i=1}^{n} \epsilon_i \leq const, \text{ for all } i = 1,\ldots,n$$

▸ $\epsilon_1,\ldots,\epsilon_n$ are *slack variables* that allow individual observations to be on the wrong side of the margin or the hyperplane

  ▸ $\epsilon_i = 0$ means observation is on the correct side of the margin, $\epsilon_i > 0$ means observation is on the wrong side of the margin. $\epsilon_i > 1$ means on the wrong side of the hyperplane

Const is a regularization parameter

▸ *Const* as a budget for the amount that the margin can be violated

- ▸ If $const = 0$, it is maximal margin hyperplane

- ▸ For $const > 0$ no more than $const$ observations can be on the wrong side of the hyperplane

- ▸ $const$ is small, we seek narrow margins that are rarely violated and vice versa

- ▸ Observations that lie directly on the margin, or *on the wrong side of the margin for their class*, are known as support vectors

Computing the Support Vector Classifier

$$\max_{\beta_0,\beta,|\beta|=1,\epsilon_0,\epsilon_1,\dots,\epsilon_n} M$$

$$y_i\left(x_i^T\beta + \beta_0\right) \geq M(1 - \epsilon_i),$$

$$\epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq const, \text{for all } i = 1, \dots, n$$

Is equivalent to
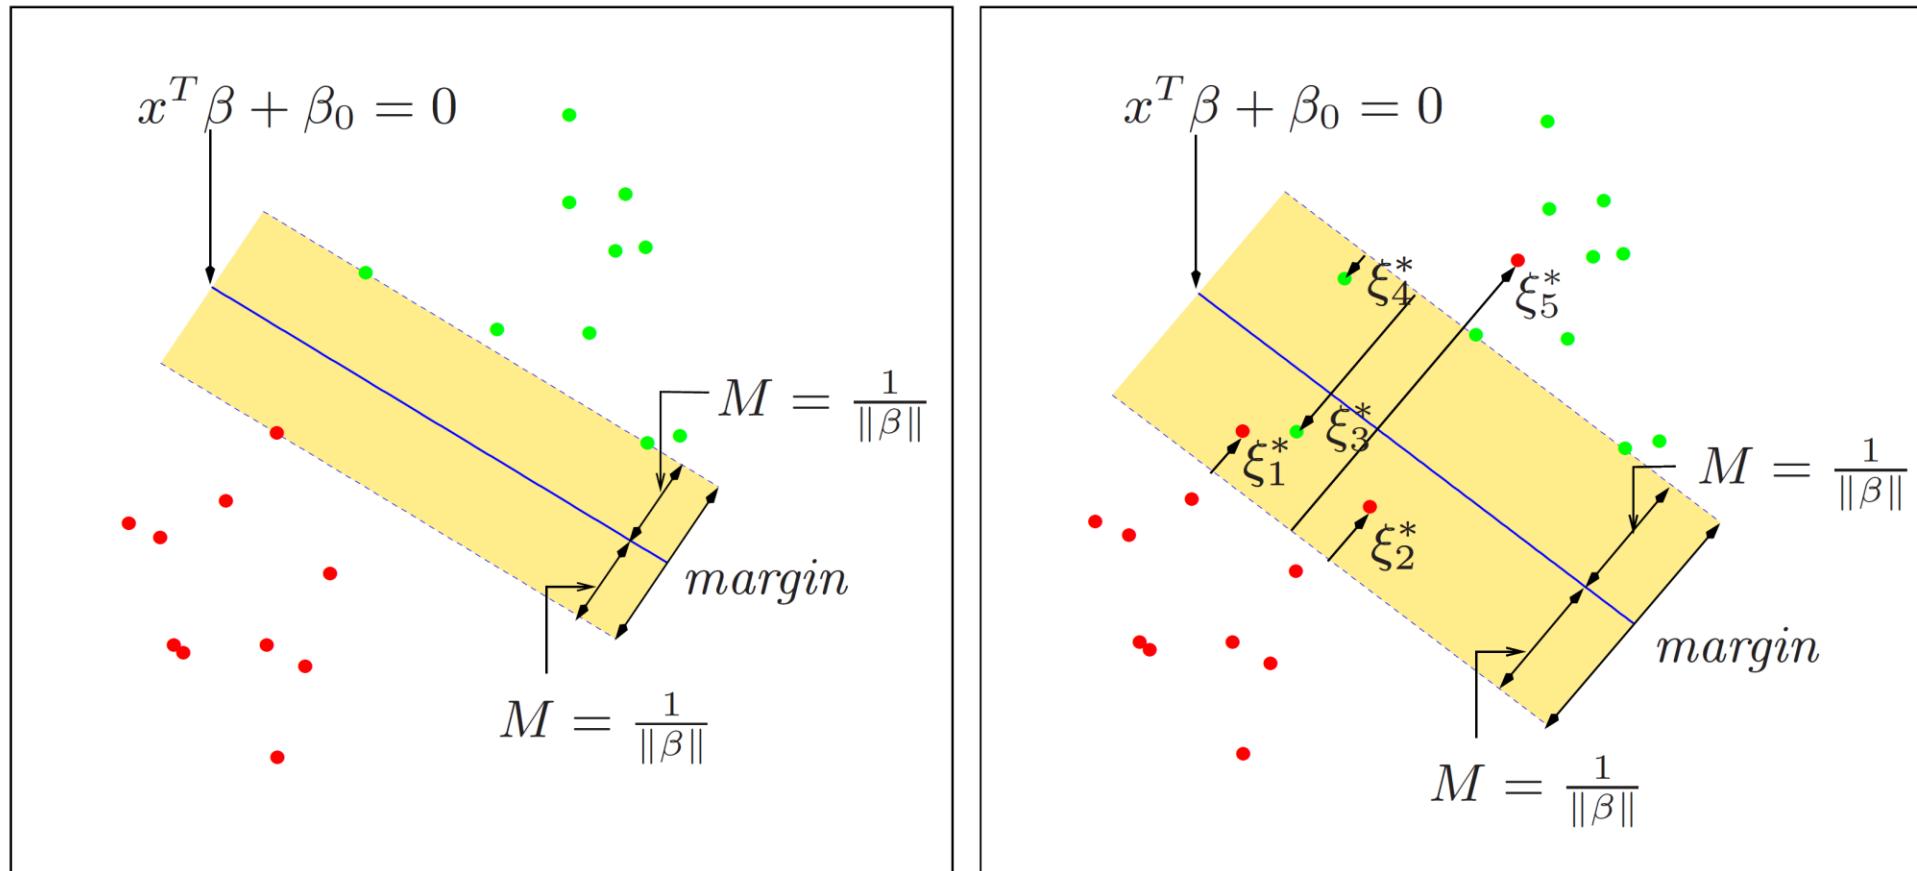
$$\max_{\beta_0,\beta,\epsilon_0,\epsilon_1,\dots,\epsilon_n} M$$

$$\frac{1}{|\beta|} y_i\left(x_i^T\beta + \beta_0\right) \geq M(1 - \epsilon_i),$$

$$\epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq const, \text{for all } i = 1, \dots, n$$

Let $M = 1/|\beta|$

$$\min_{\beta_0,\beta,\epsilon_0,\epsilon_1,\dots,\epsilon_n} |\beta|$$

$$y_i\left(x_i^T\beta + \beta_0\right) \geq (1 - \epsilon_i),$$

$$\epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq const, \text{for all } i = 1, \dots, n$$

$\sum_{i=1}^{n} \epsilon_i \leq const$

Computing the Maximal Margin Classifier (Assume $\epsilon_i = 0$)

▸ The Lagrange (primal) problem is (suitable for small p) to minimizing

$$L_P = \frac{1}{2}|\beta|^2$$

Subject to $y_i\left(x_i^T \beta + \beta_0\right) \geq 1$

$p + 1$ variables, n constraints

▸ The Lagrangian (dual) problem is (suitable for small n, can use the kernel trick) to maximizing

$$L_D = \sum_{i=1}^{n} \alpha_i - \frac{1}{2}\sum_{i=1}^{n}\sum_{i'=1}^{n} \alpha_i \alpha_{i'} \, y_i y_{i'} x_i^T x_i'$$

Subject to $0 \leq \alpha_i$ and $\sum_{i=1}^{n} \alpha_i y_i = 0$

n variables, $n + 1$ constraints

Computing the Support Vector Classifier (Assume $\epsilon_i \neq 0$)

▸ The Lagrange (primal) problem is to minimizing (C is inverse proportional to $const$)

$$L_P = \frac{1}{2}|\beta|^2 + C \sum_{i=1}^{n} \epsilon_i$$

Subject to $\epsilon_i \geq 0, y_i(x_i^T \beta + \beta_0) \geq 1 - \epsilon_i$

$p + 1 + n$ variables, $2n$ constraints

▸ The Lagrangian (dual) problem is to maximizing

$$L_D = \sum_{i=1}^{n} \alpha_i - \frac{1}{2} \sum_{i=1}^{n} \sum_{i'=1}^{n} \alpha_i \alpha_{i'} y_i y_{i'} x_i^T x_i'$$
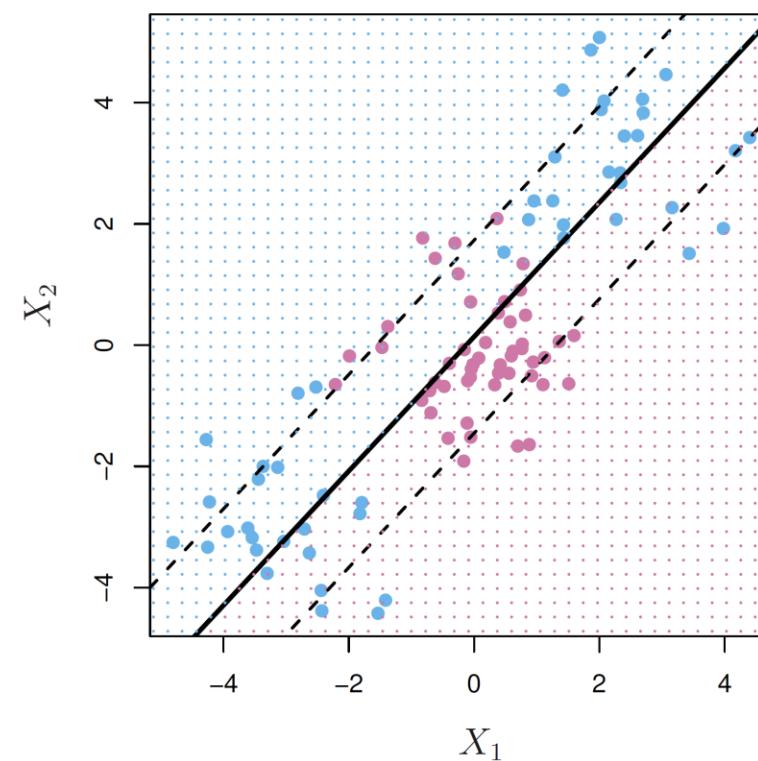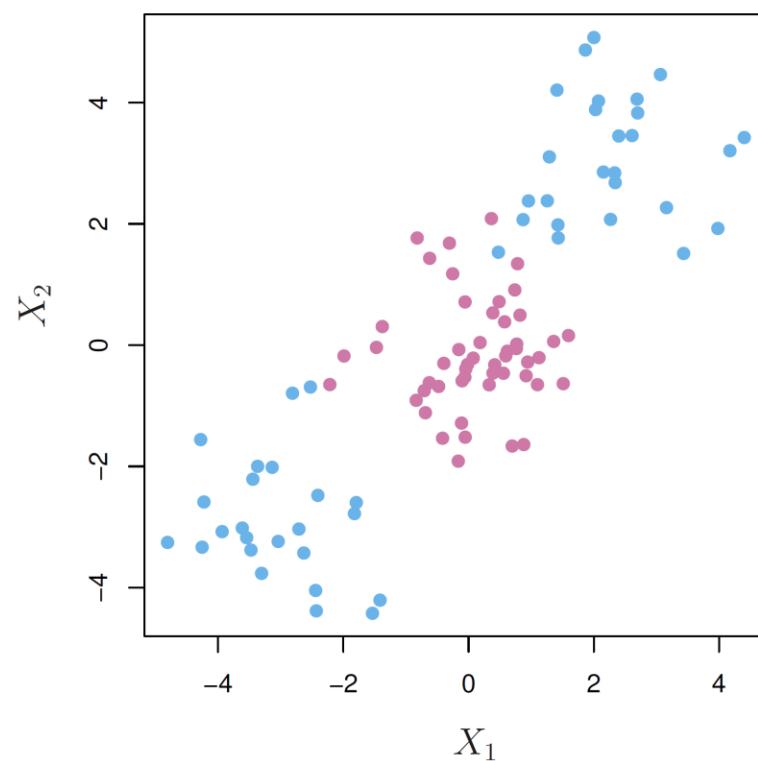
Subject to $0 \leq \alpha_i \leq C$ and $\sum_{i=1}^{n} \alpha_i y_i = 0$

n variables, $2n + 1$ constraints

▸ It is noted that the solution has the form $\hat{\beta} = \sum_{i=1}^{n} \hat{\alpha}_i y_i x_i$

Linear boundary can fail

▸ Sometimes a linear boundary simply won't work, no matter what value of *const*

▸ The example below is such a case. What to do?

Feature Expansion

▸ Enlarge the space of features by including transformations; e.g. $X_1^2, X_1^3, X_1 X_2, X_1 X_2^3, \ldots$. Hence go from a p-dimensional space to a larger dimensional space

  ▸ Fit a support-vector classifier in the enlarged space

  ▸ This results in non-linear decision boundaries in the original space

▸ Example: Suppose we use $(X_1, X_2, X_1^2, X_2^2, X_1 X_2)$ instead of just (X_1, X_2). Then the decision boundary would be of the form
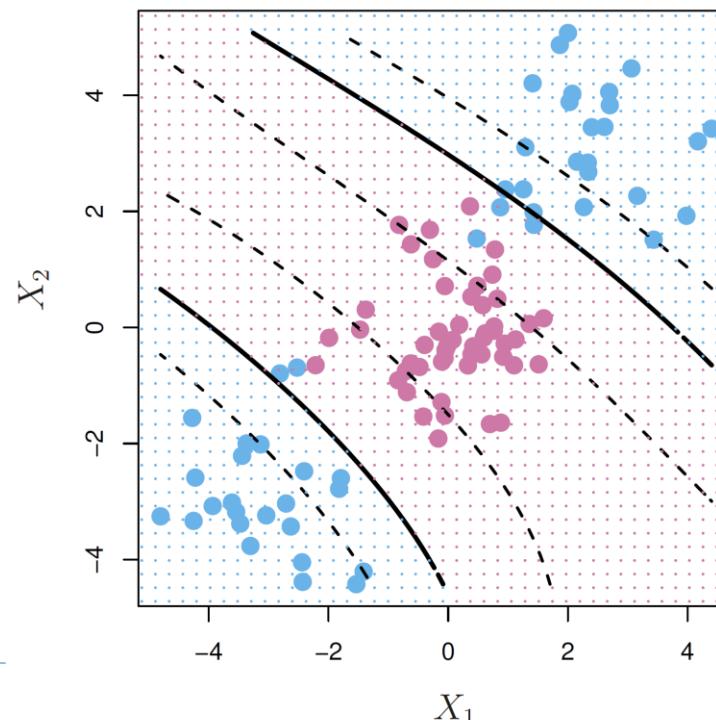
$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_1^2 + \beta_4 X_2^2 + \beta_5 X_1 X_2 = 0$$

  ▸ This leads to nonlinear decision boundaries in the original space (quadratic conic sections)

Feature Expansion - Cubic Polynomials

▸ Here we use a basis expansion of cubic polynomials. From 2 variables to 9

▸ The support-vector classifier in the enlarged space solves the problem in the lower-dimensional space

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_1^2 + \beta_4 X_2^2 + \beta_5 X_1 X_2 + \beta_6 X_1^3 + \beta_7 X_2^3 + \beta_8 X_1 X_2^2 + \beta_9 X_1^2 X_2 = 0$$

Nonlinearities and Kernels

1. Polynomials (especially high-dimensional ones) get wild rather fast
2. There is a more elegant and controlled way to introduce nonlinearities in support-vector classifiers — through the use of *kernels*

▸ Main idea: feature mapping to a high dimensional space

$$x \rightarrow \Phi(x) = (\varphi_1(x), \varphi_2(x), \varphi_3(x), \dots), \Phi(x)^T \Phi(u) = K(x, u)$$

▸ Kernel Trick: We do not really need to know $\Phi(x)$. Instead, we work on original space using the Kernel

▸ A *kernel* is a function that quantifies the similarity of two observations

▸ Recall the role of inner products in statistical learning

Inner products and support vectors

▸ Many problems involve the following inner product

$$\langle x_i, x_{i'} \rangle = \sum_{j=1}^{p} x_{ij} x_{i'j}$$

▸ The linear support vector classifier can be represented as

$$f(x) = \beta_0 + \sum_{i=1}^{n} \alpha_i y_i \langle x, x_i \rangle$$

▸ To estimate the parameters $\alpha_1, \ldots, \alpha_n$ and β_0, all we need are the $\binom{n}{2}$ inner products $\langle x_i, x_{i'} \rangle$ between all pairs of training observations

  ▸ It turns out that most of the $\hat{\alpha}_i$ can be zero except for the support vectors:

$$f(x) = \beta_0 + \sum_{i \in S} \hat{\alpha}_i y_i \langle x, x_i \rangle$$

S is the support set of indices i such that $\hat{\alpha}_i > 0$
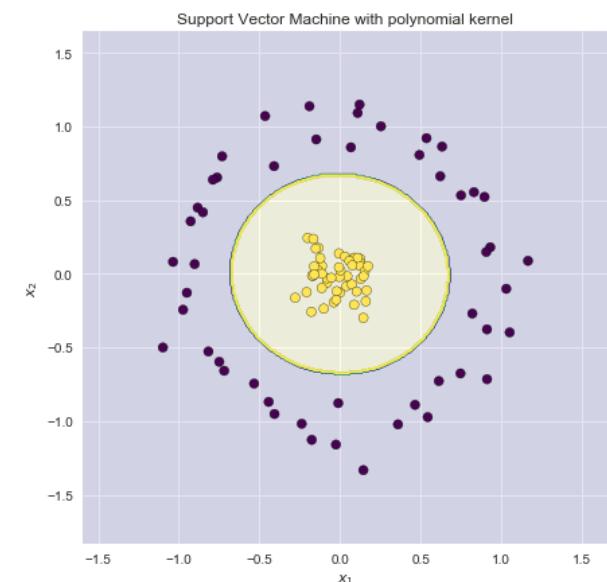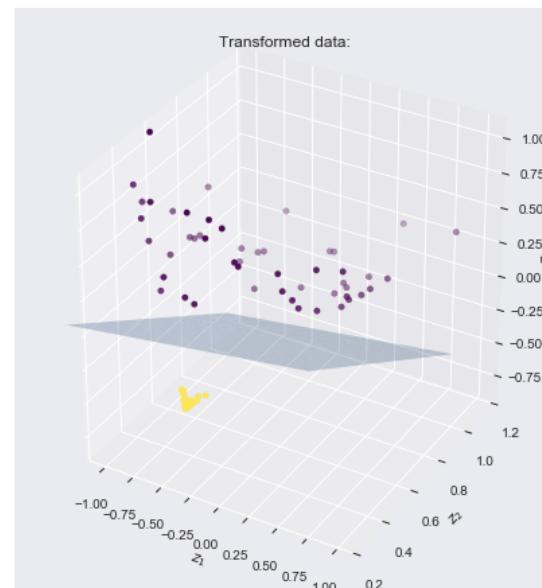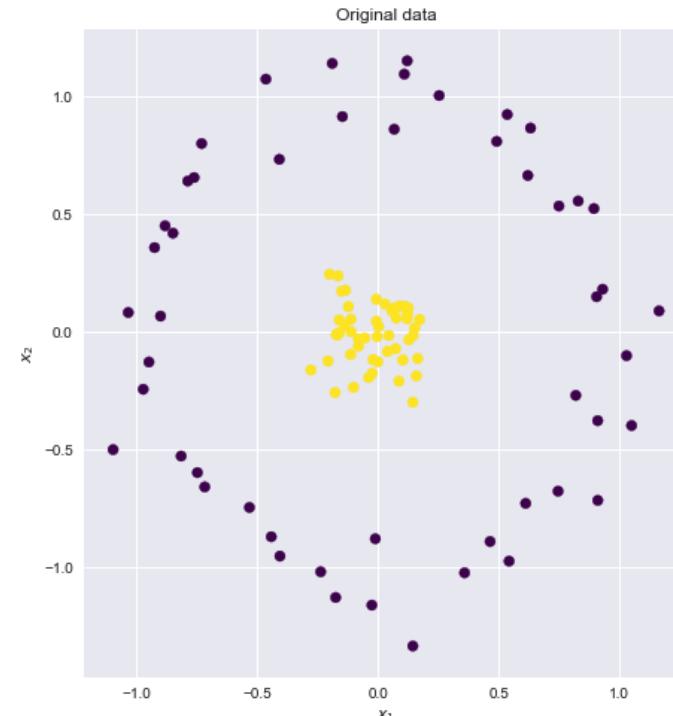
3. Kernels and Support Vector Machines

▸ The support vector machine (SVM) is an extension of the support vector classifier that results from enlarging the feature space in a specific way, using kernels

▸ Consider the second degree mapping

$$\Phi(x) = \Phi\begin{pmatrix} \sqrt{2}x_1 x_2 \\ x_1^2 \\ x_2^2 \end{pmatrix}$$

$$\Phi(a)^T \Phi(b) = \begin{pmatrix} \sqrt{2}a_1 a_2 \\ a_1^2 \\ a_2^2 \end{pmatrix}^T \begin{pmatrix} \sqrt{2}b_1 b_2 \\ b_1^2 \\ b_2^2 \end{pmatrix}$$

$$= (a_1 b_1 + a_2 b_2)^2 = \left(\begin{pmatrix} a_1 \\ a_2 \end{pmatrix}^T \begin{pmatrix} b_1 \\ b_2 \end{pmatrix} \right)^2 = (a^T b)^2$$

Original data

Transformed data:

Support Vector Machine with polynomial kernel

3. Kernels and Support Vector Machines

▸ If we can compute inner-products between observations, we can fit a SV classifier. The Support Vector Machines is an extension of SVC (linear kernel) using kernels
  In fact, SVC is a SVM with linear kernel

$$K\langle x_i, x_{i'}\rangle = \sum_{j=1}^{p} x_{ij} x_{i'j}$$

▸ Some special kernel functions are useful. E.g. the polynomial kernel
  $$K(x_i, x_{i'}) = (1 + \sum_{j=1}^{p} x_{ij} x_{i'j})^d \text{ (Add constant and variable itself)}$$

  computes the inner-products needed for d dimensional polynomials

▸ The solution now has the form

$$f(x) = \beta_0 + \sum_{i \in S} \hat{\alpha}_i y_i K\langle x, x_i\rangle$$
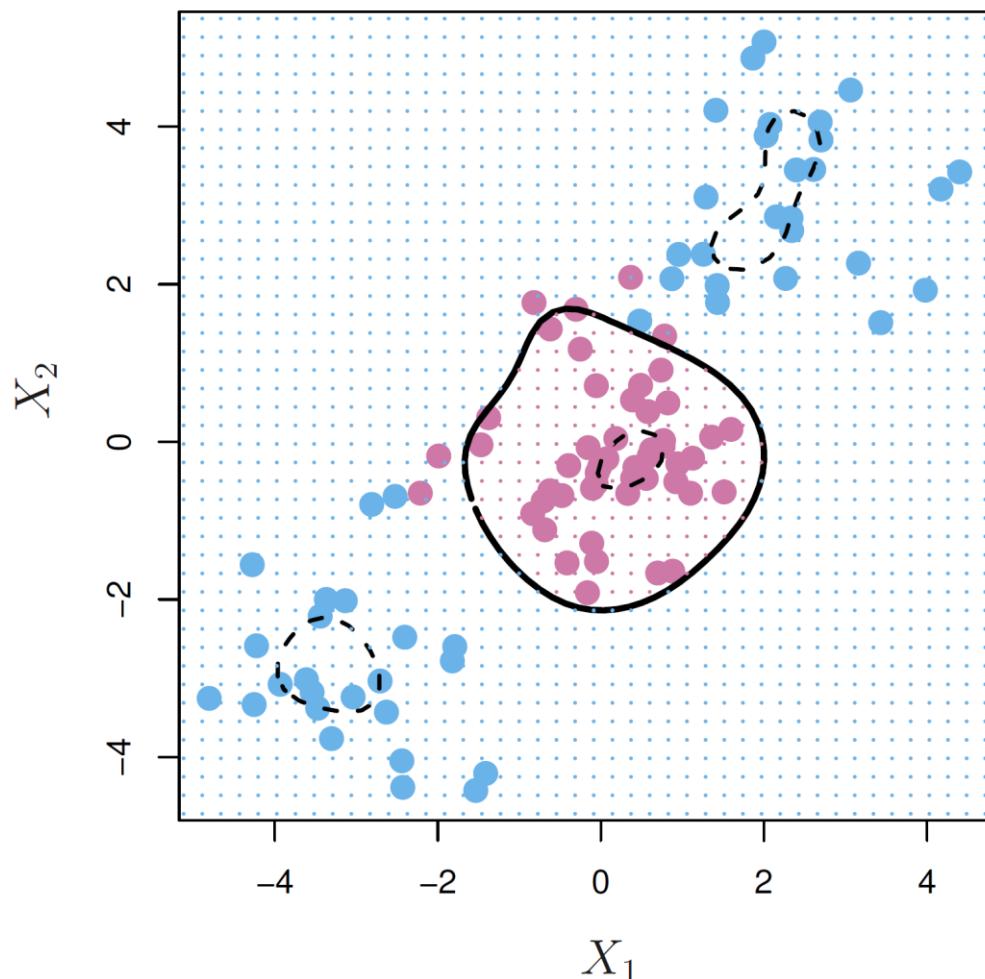
Radial Basis Kernel

▶ Controls variance by squashing down most dimensions severely

$$K\langle x, x_i \rangle = \exp(-\gamma \sum_{j=1}^{p} (x_{ij} - x_{i'j})^2)$$

$$f(x) = \beta_0 + \sum_{i \in S} \hat{\alpha}_i y_i K\langle x, x_i \rangle$$
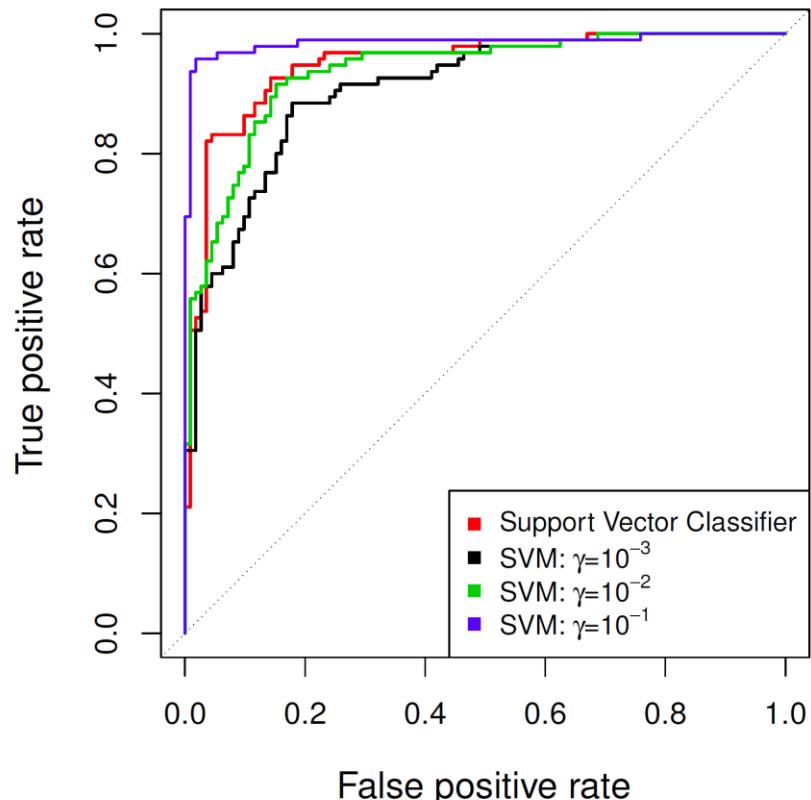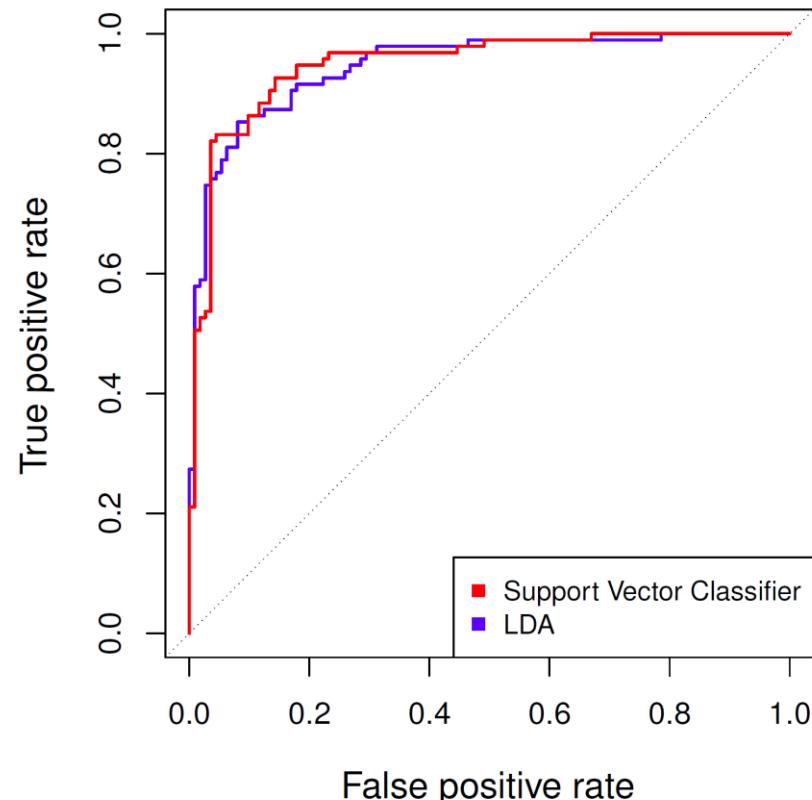
▶ γ is also a regularization parameter (You should reduce it if overfitting)

  ▶ The feature space is implicit and infinite-dimensional (By Mercer's theorem, exact Φ is intractable)
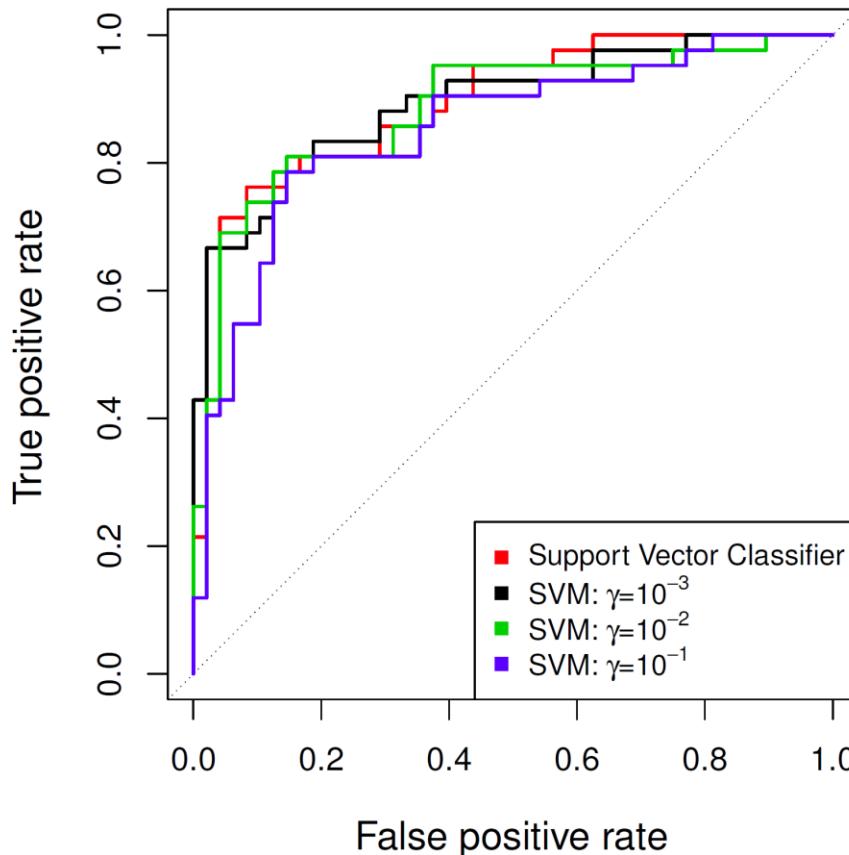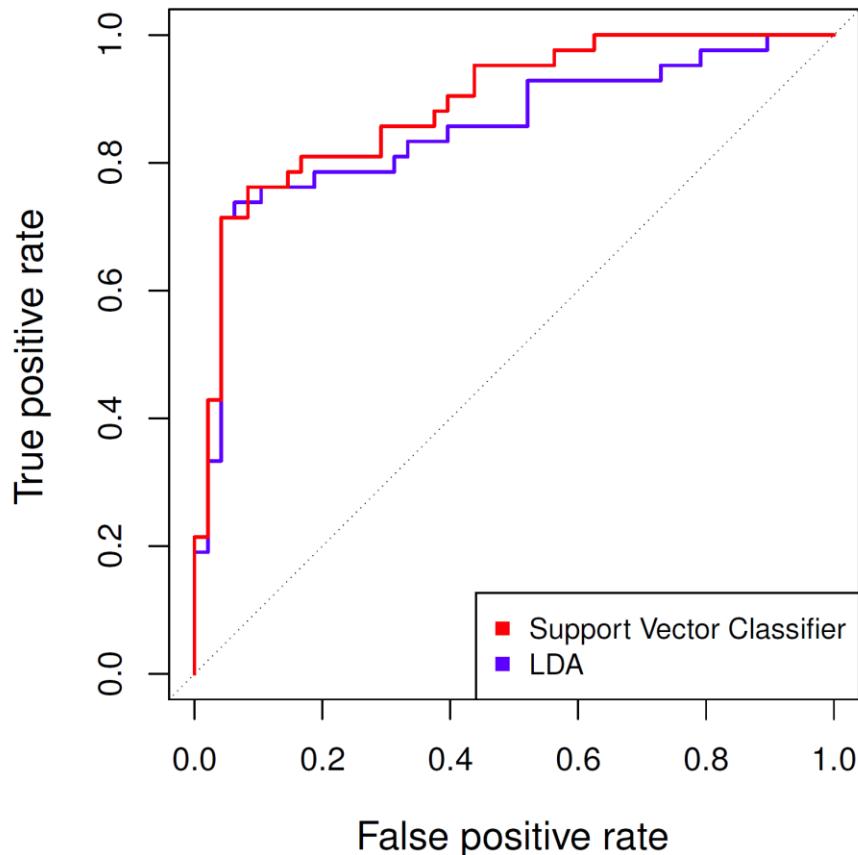
Example: Heart Data

▶ ROC curve is obtained by changing the threshold 0 to threshold t in $\hat{f}(X) \geq t$, and recording false positive and true positive rates as t varies. Here we see ROC curves on training data

Example continued: Heart Test Data

▶ ROC curves on testing data

SVMs: more than 2 classes?
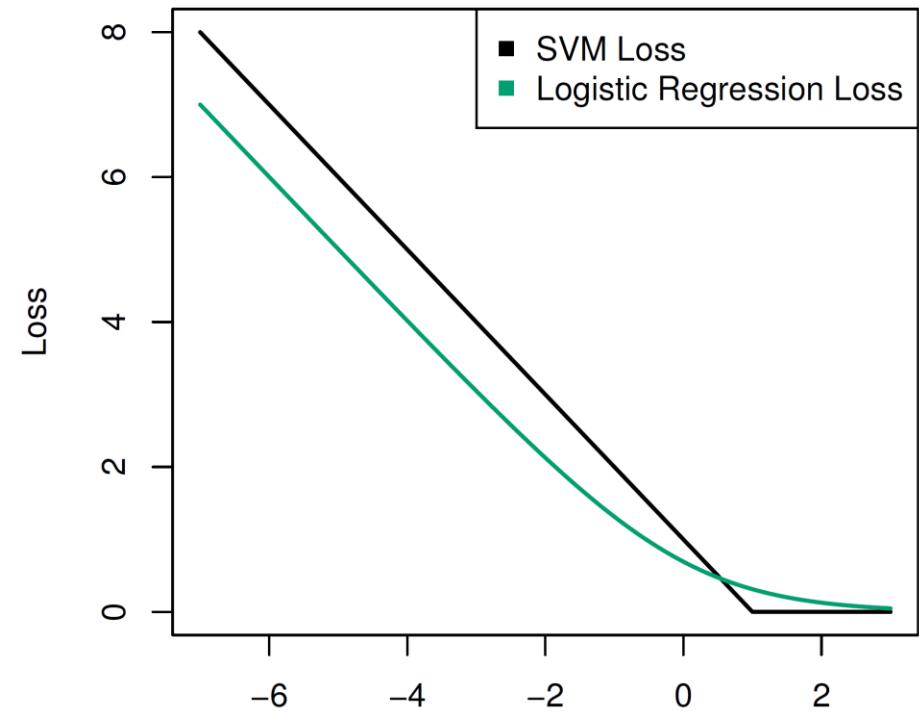
▸ The SVM as defined works for $K = 2$ classes. What do we do if we have $K > 2$ classes?

 ▸ OVA (One versus All or OVR): Fit K different 2-class SVM classifiers $\hat{f}(x),\ k = 1,\dots,K$; each class versus the rest. Classify x^* to the class for which $\hat{f}(x^*),$ is largest

  ▸ OVO (One versus One): Fit all $\binom{K}{2}$ pairwise classifiers $\hat{f}_{kl}(x)$. Classify x^* to the class that wins the most pairwise competitions

▸ Which to choose? If K is not too large, use OVO

Support Vector versus Logistic Regression?

▸ With $f(X) = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p$, we can rephrase *support-vector classifier* optimization as ([ESL exercise 12.1](#))

$$\min_{\beta_0, \beta_1, \ldots, \beta_p} \{ \sum_{i=1}^{n} \max[0, 1 - y_i f(x_i)] + \lambda \sum_{j=1}^{p} \beta_j^2 \}$$

  ▸ When λ is large then β is small, more violations to the margin are tolerated. It is proportional to *const*

  ▸ This has the form loss plus penalty. The loss is known as the *hinge loss*

  ▸ Very similar to "loss" in [logistic regression](#) (negative log-likelihood)

$y_i(\beta_0 + \beta_1 x_{i1} + \ldots + \beta_p x_{ip})$

Support Vector versus Logistic Regression?

‣ The loss function shown above is exactly zero for observations for which
$$y_i(\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p) \geq 1$$

  ‣ The loss function for logistic regression shown in Figure is not exactly zero anywhere. But it is very small for observations that are far from the decision boundary

  ‣ When classes are (nearly) separable, SVC does better than LR. So does LDA

  ‣ When not, LR (with ridge penalty) and SVC very similar

  ‣ If you wish to estimate probabilities, LR is the choice

‣ For nonlinear boundaries, kernel SVMs are popular. Can use kernels with LR and LDA as well, *but computations are more expensive*

  ‣ We could also just as well perform logistic regression or many of the other classification methods seen in this book using non-linear kernels!

SGD classifier

▸ Stochastic Gradient Descent (SGD) is a simple yet very efficient approach to fitting linear classifiers and regressors under convex loss functions such as (linear) Support Vector Machines and Logistic Regression

  ▸ Even though SGD has been around in the machine learning community for a long time, it has received a considerable amount of attention just recently in large-scale learning

  ▸ SGD is sensitive to feature scaling

▸ Kernel Approximation can be used with SGD

▸ For large-scaled problem, try to use *SGDClassifier* with hinge loss

Computations

- Support Vector Machine algorithms are <u>not scale invariant</u>, so it is highly recommended to scale your data

    - For the linear case, the algorithm used in *LinearSVC* by the *liblinear* implementation. Note that *LinearSVC* does not accept parameter kernel, as this is assumed to be linear

    - For nonlinear case, *SVC* (relies on *libsvm*) and *NuSVC* are similar methods, but accept slightly different sets of parameters

    - *liblinear and libsvm are* efficient library developed by Chih-Jen Lin

- SVM can also estimating the probability

- C is 1 by default and it's a reasonable default choice. If you have a lot of noisy observations you should decrease it: decreasing C corresponds to more regularization

SVM

▸ The key features of SVMs are the use of kernels, the absence of local minima, the sparseness of the solution and the capacity control obtained by optimizing the margin

▸ SVM can also be extended to <u>regression</u>, density estimation and novelty detection problem

  ▸ Support vector regression instead seeks coefficients that minimize a different type of loss, where only residuals larger in absolute value than some positive constant contribute to the loss function

Appendix

Reference

- https://www.csie.ntu.edu.tw/~htlin/mooc/
- https://cs229.stanford.edu/notes2021fall/cs229-notes3.pdf
  - https://cs229.stanford.edu/notes2022fall/main_notes.pdf chapter 5 and 6
- https://scikit-learn.org/stable/modules/svm.html

Statistical Modeling Y ~ X:

- *Y* continuous, *X* continuous: Regression Analysis
- *Y* discrete, *X* continuous:
  - *X* is Gaussian: Linear Discriminant Analysis
  - *X* is non-Gaussian: Logistic Regression
- *Y* discrete, *X* discrete: Discriminant Correspondent Analysis
- *Y* continuous, *X* discrete: ANOVA

Perceptron: I only need to know the hyperplane! Why bother a model?

▸ Perceptron: On-Line learning Classification

  ▸ The hyperplane is sequentially updated by the training set. Calculate the actual output and then update the weight whose increment is proportional to the difference of the actual output and the desired output

▸ Given an initial w, for each training pair (x_i, y_i) where $1 \leq i \leq n$, do the two steps:

$$o_i = f(w^T(t)x_i) \text{ (the decision rule)}$$
$$w(t+1) = w(t) + \alpha(o_i - y_i)x_i$$

From Perceptron to SVM

- Introduce the concept of Margin → a unique solution
- Introduce the slack variables → soft margin
- Introduce Kernel method → non-linear feature mapping

Common kernels

‣ Linear kernel

   ‣ Pros: safe, fast (QP solvers), explainable (by w and SVs)

   ‣ Cons: restricted (not always separable)

   ‣ A basic tool

‣ Polynomial kernel

   ‣ Pros: less restricted than linear, strong physical control by knowing the degree

   ‣ Cons: numerical difficulty for large degree

   ‣ Perhaps small degree only

‣ Gaussian kernel

   ‣ Pros: more powerful than linear and poly ones, bounded (less numerical difficulty than Poly one), only one parameter to be selected

   ‣ Cons: mysterious (no w), slower than linear, overfitting problem

   ‣ Popular but be used carefully